

KISHOR KUMAR PAROI

B.Tech / B.Sc. - Computer Science & Engineering @ RUET
Founder @ KoomAI (Live Demo: koom-ai.vercel.app | GitHub) $•$ B.Tech / B.Sc. - CSE @ RUET $•$ Dhaka, Bangladesh
1703053@student.ruet.ac.bd | kishor.ruet.cse@gmail.com | +880 1724-714156 | **koom-ai** | **macos-portfolio** | linkedin.com/in/kishor-kumar-paroi-b06344166 | github.com/KishorKumarParoi | x.com/Mr_paroi

Year	Degree / Certificate	Institute / Board	CPI / Marks	Merit Rank
2018 -- 2023	B.Sc. in Computer Science & Engineering	Rajshahi Univ. of Eng. & Tech. (RUET)	3.25 / 4.00 (D.O)	Rank 483rd in Bangladesh
2015 -- 2017	Higher Secondary Certificate (HSC)	Saint Joseph Higher Secondary School (Dhaka)	5.00 / 5.00	Rank 85th in Dhaka Board
2013 -- 2015	Secondary School Certificate (SSC)	S.J Govt Model High School (Gopalganj)	5.00 / 5.00	Rank 90th in Dhaka Board

WORK EXPERIENCE

phi0.ai *Jun'26 -- Sep'26*
Forward Deployed AI Engineer (Production RAG, GenAI & MLOps)
- Led forward deployment of **5 concurrent autonomous AI agents** orchestrating multi-step tools via **LangGraph**; achieved **99.4% task completion**, accelerating client onboarding by **4x**.
- Architected high-throughput **Hybrid RAG** with HNSW, BM25, and multi-stage re-ranking (**FlashRank + Cohere v3.5**) under **<35ms semantic cache**; cut hallucinations by **46%** and saved **$18K/mo**.
- Built continuous **MLOps & GenAI flywheel** with NeMo Guardrails and Canary K8s deployments; spearheaded enterprise AI transformation, securing **$1.4M in client contracts**.

Samsung AI Core / Care Health *Jan'26 -- May'26*
Data Engineering & AI Systems Intern
- Built batch & streaming pipelines via **AWS S3, EMR, Athena, PySpark**; cut operational costs by **$4.2K/run (34% savings)** through event-driven Lambda/SQS dispatch.
- Accelerated 8B/70B model inference on H100 GPUs via **FlashAttention-2** and fused kernels; achieved **3.8x throughput** on >32K context windows, cutting p99 latency by **54%**.
- Deployed real-time health insurance claim fraud detection pipeline over **AWS Kinesis and OpenTelemetry**; slashed false-positive claim investigations by **52%**, saving **$380K annually**.

kepi.ai *Jan'22 -- Sep'25*
Senior Full Stack Engineer (3 yrs 9 mos, Remote)
- Architected microservices migration to **Kubernetes on AWS**; scaled platform capacity by **300%** while cutting p95 latency by **48%** at 99.99% availability.
- Spearheaded enterprise AI initiatives: delivered autonomous multi-agent pipelines and scalable LLM microservices, driving a **22% increase in customer LTV**.
- Hardened zero-trust mesh infrastructure with **Istio mTLS and fine-grained RBAC**; eliminated unencrypted service-to-service transit and reduced cross-zone egress costs by **38%**.

FOUNDER VENTURE: KOOMAI

KoomAI — High-Performance Loom Alternative Live *Jan'25 -- Dec'25*
Founder & Principal Architect
- **WebRTC Low-Latency Engine:** Engineered real-time WebRTC media pipeline (**<120ms** latency) with adaptive bitrate streaming (SVC); sustained 60 FPS 1080p video with 0% stutter across **50K+ active video minutes**.
- **ElectronJS Desktop Client:** Built cross-platform desktop application leveraging native OS-level IPC audio/display capture, hardware-accelerated NVENC/VideoToolbox encoding; reduced idle CPU usage by **45%**.
- **Chrome Extension (Manifest V3):** Developed lightweight screen-recording browser extension utilizing offscreen documents for synchronized tab/mic audio mixing, capturing **10K+ user recordings**.
- **Next.js 16 Cockpit & Player:** Created responsive video cockpit using **Next.js 16 App Router**, React Server Components, and waveform scrubbing; cut initial page load (FCP) by **58%**.
- **High-Throughput Express API & CDN:** Architected ingestion backend on Express.js with **AWS S3 multipart uploads** and CloudFront edge delivery; slashed video storage operational costs by **62%**.
- **Monetization & Enterprise Workspaces:** Integrated end-to-end **Stripe billing engine** with idempotent webhook handling and multi-tenant team access controls, driving **$24K MRR** in pilot adoption.

ACHIEVEMENTS & HACKATHONS

- **Competitive Programming & DSA: Ranked 232nd** in Codeforces Div-2 (**Blue Coder**), **5★** in CodeChef, **17th in IUPC**; Qualified **Meta Hacker Cup R3**; 1000+ DSA on LeetCode.
- **Champion | Technobeast AI Agents Hackathon:** Engineered Qdrant vector optimization with negative embeddings and trajectory matching.
- **7th Rank | AI Security Hackathon:** Ranked in top tier for real-time prompt injection firewall (**97.2% accuracy**, **<30ms** latency).
- **National Merit Scholarships:** Conferred merit scholarships for 85th Rank in Dhaka Board (HSC) and 90th Rank in Dhaka Board (1st in Gopalganj District) (SSC).

CERTIFICATIONS & KEY COURSES

- **DeepLearning.AI / Stanford (Andrew Ng):** ML & Deep Learning Specializations (Neural Nets, CNNs, Transformers, Vector Embeddings).
- **Distributed Systems & AI Architecture:** High-Throughput Systems & AI Architecture (Arpit Bhayani, Gaurav Sen, ByteByteGo), MLOps (Krish Naik), LeetCode DSA.
- **Core CS Coursework (RUET):** Data Structures & Algorithms, Operating Systems, DBMS, Distributed Systems, Computer Networks.

LEADERSHIP & MENTORING

- **In-Person AI Agents Masterclass:** Conducted hands-on in-person AI Agents building workshops on **LangGraph state machines**, tool calling, and production multi-agent deployment.
- **Technical Mentorship:** Mentored **100+ students & engineers** across Machine Learning, AI Agents, DevOps, and MLOps, guiding end-to-end cloud production pipelines.

TECHNICAL PROJECTS

One-Click Deploy: curl | bash

Zomato AI Lakehouse & Multimodal Platform Deploy *Sept'26 -- Present*
Lakehouse (35M+ Rows), Feast Store & CV Engine
- Architected enterprise Lakehouse on **Snowflake & PySpark Delta Lake** processing **35M+ orders (2.3 GB)** with **dbt Core SCD-2**; cut warehouse compute by **42% ($18K/mo savings)**.
- Engineered real-time feature streaming via **Airflow**, **Debezium CDC WAL**, and **Feast Store** over Redis; cut serving latency to **<5ms**, driving **+$340K quarterly revenue lift**.
- Deployed multimodal CV dispute arbitration (**dHash** anti-fraud, receipt OCR) with **G-Eval CoT**; slashed fraudulent claims by **68%**, protecting **$1.2M in annual partner margins**.

Hybrid Anime Intelligence & LLMOps Platform Deploy *Aug'26 -- Present*
Polyglot Mesh (Go/Rust), RSLoRA Flywheel & Multi-LoRA vLLM
- Polyglot mesh: **Go 1.24** auth, **Rust 1.85 SIMD** scoring, and **FastAPI** AI core over **Kafka**; cut scoring latency from 25ms to **<0.5ms (50x speedup)**, saving **64% CPU cost**.
- Engineered hybrid retrieval combining **Dual-Tower NCF**, **ChromaDB HNSW**, and **BM25** via Reciprocal Rank Fusion (**NDCG@10 = 0.94**); drove a **28% lift in user session retention**.
- Served adapters via **vLLM Multi-LoRA dynamic multiplexing** on 1x GPU with **RSLoRA flywheel**; slashed GPU spend by **68%** while scaling to **3.4x higher token throughput**.

Enterprise AI Security & Agentic RAG Platform Deploy *Oct'26 -- Present*
Defense-in-Depth AI Gateway, NVIDIA NeMo & LangGraph
- Built zero-trust defense with **BotShield WAF** and **NVIDIA NeMo Guardrails**; blocked **95%+ of prompt injections & jailbreaks** (benchmarked via **Microsoft PyRIT 0.14.0**).
- Engineered stateful agentic core in **LangGraph** with MemorySaver checkpoints and **Portkey Gateway**; saved **$8.4K/month (42%)** via **<50ms semantic similarity caching**.
- Standardized tool execution via **FastMCP** with **Human-in-the-Loop (HITL)** gates and Logfire/LangSmith tracing; achieved **100% audit compliance** with zero unauthorized tool runs.

Smart Manufacturing Industrial AI & MLOps Deploy *Oct'26 -- Present*
Polyglot (Rust/Go/Next.js), PyTorch Autoencoders, GKE/K3s & ArgoCD
- Deployed IoT anomaly detection (**PyTorch Autoencoders**) and defect inspection (**CNN**, <20ms); reduced part escape defect rate by **31%**, saving **$420K annually in scrap waste**.
- Built polyglot edge architecture: **Rust 1.85 SIMD** vibration pre-filter (:8082), **Go cloud gateway** (:8085), **Next.js 15** cockpit (:3000), and **PyTorch AI core** (:8000).
- Automated declarative CI/CD via **Jenkins, SonarQube, Nexus 3, and ArgoCD GitOps**; cut deployment cycle by **96% (4h to 8m)** with zero-downtime failover.

Core Systems: Rust Git Engine & POSIX Terminal Shell *Aug'24*
Bare-Metal Systems in Rust 1.85, Concurrency & Pipelines
- Engineered Git internals in **Rust 1.85** (SHA-1 object graph, streaming zlib SIMD decompression); matched C speed with **100% compile memory safety**.
- Built async POSIX terminal shell with signal multiplexing and I/O redirection; achieved **sub-μs latency** with **73% lower heap allocation**.

TECHNICAL SKILLS

Languages & Frameworks: Python, C/C++, Rust (Actix), Go (Gin), SQL, Bash, TypeScript, Java (Spring Boot), FastAPI, Express, NestJS, Next.js 16, Streamlit
AI, ML & LLMOps: PyTorch, vLLM, LangGraph, LangChain, HuggingFace, TRL, PEFT (RSLoRA, QLoRA), NeFTune, MLflow, Feast, Kubeflow, AWS SageMaker, KServe, DVC
Big Data & Lakehouse: Snowflake, Databricks, Azure Data Factory (ADF), Synapse CETAS, ADLS Gen2, dbt Core, PySpark, Delta Lake, Kafka, Airflow, Debezium CDC
Cloud Platforms & Infra: AWS, GCP, Azure, Terraform, Ansible, Helm, ArgoCD, Sonatype Nexus, Docker, Kubernetes (GKE, EKS, K3s)
DevSecOps & Observability: NVIDIA NeMo Guardrails, PyRIT, BotShield WAF, SonarQube, Trivy, Jenkins, GitLab CI/CD, Prometheus, Grafana, EFK (Elasticsearch/Kibana)
WebRTC, 3D & Full-Stack: WebVR, WebRTC, Three.js, React Three Fiber, ElectronJS, Next.js 16, Prisma, Drizzle, Stripe Connect, PostgreSQL, MongoDB Atlas